

Is this a good text instruction: "When filtering by company name, do not use wildcard matching (LIKE %name%). Instead match the user's requested entity against the exact, full official name found in the Name column using an exact string match (=). If a user provides an incomplete or generic company name that maps to multiple potential records, do not guess or execute an arbitrary lookup. Actively generate a clarification response prompting the user to select the specific company they want."

Weighted instruction design against entity matching and SQL best practices

This is a mixed bag — the second half is a strong example of a good text instruction, but the first half is fighting the wrong tool.

The problem: the wildcard/exact-match part

Telling Genie "don't use LIKE, use =" is dictating exact SQL syntax — which the docs explicitly warn against. General instructions work best as conceptual guidance ("this is what the data means, this is the domain"), not as SQL-writing rules. There are two issues with embedding it as text:

1. It's **unreliable**. Genie has to correctly recall and apply this rule every time it generates a query, and compliance with syntax-level text rules tends to degrade in longer conversations or as instructions compete with other guidance.
2. There's a **purpose-built feature for this**. Databricks has an **entity matching / prompt matching** setting you can enable directly on the **Name** column (Configure > Data > column config > entity matching). When enabled, Genie matches user input against the actual stored values in that column — correcting spelling and mapping loosely-phrased input to the exact stored name — and the filter renders as a dropdown of real values rather than an LLM-generated **LIKE** guess. That gets you exact-match behavior structurally, without relying on the model. ↓ remember a rule.

Suggested revision:

1. Enable **entity matching** on the **Name** column (handles the exact-match/no-wildcard behavior structurally).
2. Keep a trimmed text instruction focused only on the disambiguation behavior:

"When a user's company reference matches multiple records in the Name column, do not guess. Ask the user to clarify which specific company they mean before proceeding."

following the above instruction are these two:

1) anytime a user is asking for the company SigFood do not use a Wildcard in the SQL Statement like Name = %sigfood%. Instead use entity matching to filter for SIG FOOD INC

2) if entity matching on company name cannot identify a unique company provide the user a list of companies that could match their request and ask them whether they want to be more specific or proceed running the query with all the matches.

Are these three questions overlapping or addressing the same underlying issue? Or are they good as they are? what would you recommend for the best genie space curation?

Identified overlapping instructions and recommended consolidating synonym mapping with unified mat...

Good catch bringing these together — yes, there's meaningful overlap, and one of these shouldn't be a text instruction at all.

Instruction 1 (SigFood → SIG FOOD INC): this shouldn't be a text instruction

This is a single hardcoded name mapping, and the docs are pretty direct about this anti-pattern: hardcoded specific values belong in the data layer or in configuration, not in general text instructions. Two problems:

- **It doesn't scale.** If you have dozens of nicknames/abbreviations across your company list ("SigFood," "Acme Corp" vs "ACME Corporation," etc.), you'd need one instruction per company — burning through your instruction budget fast and creating a wall of text Genie has to parse and prioritize on every query.
- **There's already a feature for exactly this.** The synonyms setting on the **Name** column (part of entity matching config) is built to map informal/partial names to the canonical stored value. Add "SigFood" as a synonym for "SIG FOOD INC" there, and entity matching resolves it structurally — no text instruction needed, and it's more reliable than hoping the LLM recalls a prose rule.

Recommendation: Delete this as a text instruction. Move it into the column's synonym list instead. Repeat for any other known aliases you're aware of.

Instruction 2: good behavior, but it overlaps with (and should replace) the clarification instruction from last time

This is well-formed — clear trigger ("entity matching cannot identify a unique company") and clear expected action. But it's addressing the *same scenario* as the instruction we landed on previously ("ask the user to clarify which specific company they mean"), just with a more refined resolution path (offer the list, let the user choose to narrow or proceed with all matches).

Keeping both is a problem: they're two instructions governing the same trigger condition with different resolution logic, which creates ambiguity about which one Genie should actually follow, and eats two instruction slots for one behavior.

Recommendation: Replace the earlier "ask to clarify" instruction entirely with this one — don't stack them.

Consolidated curation plan

Structural (knowledge store / column config, not text):

- Enable entity matching on `Name`
- Add synonyms for known aliases (SigFood → SIG FOOD INC, and others as you find them)

Single text instruction (replaces prior version):

"When a user's company reference does not map to a unique record via entity matching, present the list of potential matching companies and ask whether they want to specify one or proceed with the query using all matches."

This gets you: reliable name resolution handled by configuration (scales to any number of aliases without bloating instructions), and one unambiguous behavioral rule for the ambiguous case — rather than three overlapping instructions competing for priority.